

Dictionary data structures

William Hendrix

Outline


- Review
- Other dictionary implementations
 - Binary search trees
 - Hash tables
 - Bit vectors
- Heaps
- Union-find
- Prefix and suffix trees

Dictionary

- Abstract data structure for storing and retrieving values
- **Primary operations**
 - *Search(x)*: returns the location of x in the dictionary, or NIL if not contained
 - *Insert(x)*: adds x to the dictionary
 - *Delete(x)*: removes x from the dictionary
- **Additional operations**
 - *Max()*, *Min()*: return the location of the largest/smallest element
 - *Successor(x)*, *Predecessor(x)*: return the next largest/smallest element than x

Linked list operations

- Search: linear scan




- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain




Linked list operations

- Search: linear scan




- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain




Linked list operations

- Search: linear scan




- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain




Linked list operations


- Search: linear scan



- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain




- If sorted, scan to find insertion position
- Delete: reroute links, then delete victim




Linked list operations


- Search: linear scan



- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain




- If sorted, scan to find insertion position
- Delete: reroute links, then delete victim




Linked list operations


- Search: linear scan



- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain




- If sorted, scan to find insertion position
- Delete: reroute links, then delete victim




Linked list operations

- Search: linear scan



- If sorted, stop when values are too large
 - If DLL, you can search backwards from end or forwards
- Insert: add links to include new node in chain



- If sorted, scan to find insertion position
- Delete: reroute links, then delete victim



- If SLL, need to scan to find previous node

Summary: link-based dictionaries

Operation	Unsorted SLL	Unsorted DLL	Sorted SLL	Sorted DLL
Search(x)	$O(n)$	$O(n)$	$O(n)$	$O(n)$
Delete(x)	$O(n)$	$O(1)$	$O(n)$	$O(1)$
Insert(x)	$O(1)$	$O(1)$	$O(n)$	$O(n)$
Build	n/a	n/a	$O(n \lg n)$	$O(n \lg n)$
Min()	$O(n)$	$O(n)$	$O(1)$	$O(1)$
Max()	$O(n)$	$O(n)$	$O(1)$	$O(1)$
Predecessor(x)	$O(n)$	$O(n)$	$O(n)$	$O(1)$
Successor(x)	$O(n)$	$O(n)$	$O(1)$	$O(1)$

- **Note:** DLL time is strictly better, asymptotically
 - Trade-off: more space, more pointer manipulation


Summary: linear dictionaries

Operation	Unsorted array	Unsorted DLL	Sorted array	Sorted DLL
Search(x)	$O(n)$	$O(n)$	$O(\lg n)$	$O(n)$
Delete(x)	$O(1)$	$O(1)$	$O(n)$	$O(1)$
Insert(x)	$O(1)$, <i>amortized</i>	$O(1)$	$O(n)$	$O(n)$
Build	n/a	n/a	$O(n \lg n)$	$O(n \lg n)$
Min()	$O(n)$	$O(n)$	$O(1)$	$O(1)$
Max()	$O(n)$	$O(n)$	$O(1)$	$O(1)$
Predecessor(x)	$O(n)$	$O(n)$	$O(1)$	$O(1)$
Successor(x)	$O(n)$	$O(n)$	$O(1)$	$O(1)$


- Arrays are usually preferred, due to lower coefficients

Binary search trees

- Non-linear linked data structure
- Trees start with a *root* node
 - Usually depicted at top
- Each node has two children
 - Use NIL link if no child on left/right
- Nodes also generally store *parent* pointer
 - NIL for root
- **Binary Search Tree Property**
 - All children of left child are equal or smaller
 - All children in right child are equal or larger




Binary tree lingo



- *Leaf*: node with no children
- *Level*: number of links away from the root
- *Height*: the max level in the tree
- *Complete*: every node above last level has two children
- *Left/right subtree* of a node: tree rooted at node's left/right child
 - Trees are *recursive* data structures
- *Balanced*: each node's left and right subtrees are similar size
- *Degenerate*: BST with only left or right children


BST properties

Complete BSTs



- Level i has 2^i nodes
- Total nodes: $\sum_{i=0}^h 2^i = 2^{h+1} - 1$
- Height: $O(\lg n)$
- Leaves: approx. $n / 2$
- Root contains median element

Degenerate BSTs



- Every level has 1 node
- Total nodes: n
- Height: $O(n)$
- Leaves: 1
- Essentially a sorted linked list
 - Left children: descending
 - Right children: ascending

BST dictionary operations

- **Search(x)**

- Binary search
- Start with Search(root, x)
- $O(1)$ time per call
- *Worst case: h calls (height of BST)*
- $O(h)$ time

```
1 Algorithm: Search(node, x)
2 if node = NIL then
3   | return NIL;
4 else if node.data = x then
5   | return node;
6 else if node.data > x then
7   | return Search(node.left, x);
8 else
9   | return Search(node.right, x);
10 end
```

- **Insert(x)**

- Binary search
- Insert as root or call Insert(root, x)
- $O(h)$ time

```
1 Algorithm: Insert(node, x)
2 if  $x \leq$  node.data then
3   | if left = NIL then
4     | left = NewNode(x);
5   else
6     | Insert(left, x);
7   end
8 else
9   | if right = NIL then
10    | right = NewNode(x);
11  else
12    | Insert(right, x);
13  end
14 end
```


BST dictionary operations

- **Delete(x)**
 - Binary search
 - Special cases depending on children
 - 0 children: delete
 - 1 child: replace w/ child
 - 2 children: find right ST min, swap
 - Worst case analysis
 - $O(h)$ to find x
 - $O(h)$ to find RST min
 - $O(h)$ time

```
1 Algorithm: Delete(node, x)
2 if node = NIL then
3   | return;
4 else if node.data > x then
5   | Delete(node.left, x);
6 else if node.data < x then
7   | Delete(node.right, x);
8 else
9   | if node.left = NIL and node.right = NIL then
10  |   | Set node.parent's child pointer to NIL;
11  |   | free node;
12  | else if node.left ≠ NIL and node.right ≠ NIL
13  |   | then
14  |     | sub = min(node.right);
15  |     | Remove sub.parent's child link;
16  |     | Set sub's 3 links to match node;
17  |     | Set sub's parent's and child's links to sub;
18  |     | free node;
19  | else if node.left ≠ NIL then
20  |   | Set parent's child pointer to node.right;
21  |   | node.right.parent = node.parent;
22  |   | free node;
23  | else
24  |   | Set parent's child pointer to node.left;
25  |   | node.left.parent = node.parent;
26  |   | free node;
27 end
```

Balanced Binary Search Trees

- How tall are BSTs?
 - Best case: $O(\lg n)$
 - Average case: $O(\lg n)$
 - Worst case: $O(n)$
- Balanced BSTs
 - Sophisticated variants of BST
 - *Guarantee* $O(\lg n)$ height with constant overhead
 - Red-Black trees, AVL trees, etc.
 - We are not going to cover details of Balanced BSTs



Summary: BST dictionaries

Operation	Binary Search Tree	Balanced BST	Unsorted array	Sorted array
Search(x)	$O(h)$	$O(\lg n)$	$O(n)$	$O(\lg n)$
Delete(x)	$O(h)$	$O(\lg n)$	$O(1)$	$O(n)$
Insert(x)	$O(h)$	$O(\lg n)$	$O(1)^*$	$O(n)$
Build	$O(n \lg n)$	$O(n \lg n)$	n/a	$O(n \lg n)$
Min()	$O(h)$	$O(\lg n)$	$O(n)$	$O(1)$
Max()	$O(h)$	$O(\lg n)$	$O(n)$	$O(1)$
Predecessor(x)	$O(h)$	$O(\lg n)$	$O(n)$	$O(1)$
Successor(x)	$O(h)$	$O(\lg n)$	$O(n)$	$O(1)$

- **Advantage:** $O(\lg n)$ is *much* better than $O(n)$ for large data
- **Disadvantage:** $O()$ hides larger coefficients for BSTs

Hash tables

- Sparse array-based data structure
- Insert elements according to a *hash function*
 - Function that maps elements in domain to integers 0 to size of array minus one ($m-1$)
 - Must take $O(1)$ time
- **Example hash function**
 - $f : \mathbb{Z} \rightarrow [0, m - 1]$
 - $f(x) = x \bmod m$
 - Most hash functions use modulus to ensure range
- **Hash table example**
 - Size = 10, hash function: mod 10
 - Inserting 3, 15, 27, 82, 96, 100

100		82	3		15	96	27		
-----	--	----	---	--	----	----	----	--	--

Collisions

- What do we do when two values map to the same location?

100		82	3		15	96	27		
-----	--	----	---	--	----	----	----	--	--

↑
25

- Insert 25
- Two basic solutions
- Separate chaining
 - Each location is the head of a linked list
 - Append new element to list
 - **Never "need" to reallocate**
- Open addressing
 - Find the next open location, insert there
 - Can scan quadratically to avoid "congestion"
 - **No links, so table can be larger with same memory**
 - **Deleting an element requires reinserting everything that follows**
- Both potentially require scanning to find element

Operations

- **Search(x)**
 - Hash element
 - Scan linked list (or until empty location)
 - Worst case: $O(n)$
- **Insert(x)**
 - Hash element
 - Append to linked list (or scan for open location)
 - Worst case: $O(1)$ (or $O(n)$)
- **Delete(x)**
 - Hash element
 - Delete from linked list (or scan/delete/re-insert)
 - Worst case: $O(n)$ (or $O(n^2)$)

```
1 Algorithm: Search(x)
2 loc = Hash(x);
3 return table[loc].Search(x);
```

```
1 Algorithm: Insert(x)
2 ins = NewNode(x);
3 loc = Hash(x);
4 ins.next = table[loc];
5 table[loc] = ins;
```

```
1 Algorithm: Delete(x)
2 loc = Hash(x);
3 node = table[loc];
4 if node.value = x then
5 |   table[loc] = node.next;
6 |   free node;
7 else
8 |   while node.next ≠ NIL do
9 |       |   next = node.next;
10 |       |   if next.value = x then
11 |           |   node.next = next.next;
12 |           |   free next;
13 |           |   node = node.next;
14 |       end
15 end
```

Hash table complexity

Operation	Separate chaining	Open addressing	Balanced BST
Search(x)	$O(n)$	$O(n)$	$O(\lg n)$
Delete(x)	$O(n)$	$O(n^2)$	$O(\lg n)$
Insert(x)	$O(n)$	$O(n)$	$O(\lg n)$
Build	$O(m + n^2)$	$O(m + n^2)$	$O(n \lg n)$
Resize	$O(m + n^2)$	$O(m + n^2)$	n/a
Min()	$O(m + n)$	$O(m)$	$O(\lg n)$
Max()	$O(m + n)$	$O(m)$	$O(\lg n)$
Predecessor(x)	$O(m + n)$	$O(m)$	$O(\lg n)$
Successor(x)	$O(m + n)$	$O(m)$	$O(\lg n)$

- This is awful!
- Why would anyone ever use a hash table?

Why would anyone use a hash table?

- Bad worst-case complexity but great *expected-case* complexity
- Expected-case assumptions
 - Hash function produces $O(1)$ collisions
 - Each inserted value has $O(1)$ duplicates
 - $m = O(n)$
- Search(x)
 - Hashing and scanning take $O(1)$ time
- Insert(x)
 - Hashing and scanning take $O(1)$ time
- Delete(x)
 - Hashing and scanning take $O(1)$ time
 - Reinsertion takes $O(1)$ time (open addressing)

Expected-case complexity

Operation	Separate chaining	Open addressing	Balanced BST
Search(x)	$O(1)$	$O(1)$	$O(\lg n)$
Delete(x)	$O(1)$	$O(1)$	$O(\lg n)$
Insert(x)	$O(1)$	$O(1)$	$O(\lg n)$
Build	$O(n)$	$O(n)$	$O(n \lg n)$
Resize	$O(1)$, amortized	$O(1)$, amortized	n/a
Min()	$O(n)$	$O(n)$	$O(\lg n)$
Max()	$O(n)$	$O(n)$	$O(\lg n)$
Predecessor(x)	$O(n)$	$O(n)$	$O(\lg n)$
Successor(x)	$O(n)$	$O(n)$	$O(\lg n)$

- This is amazing!
- The three most important techniques are hashing, hashing, and hashing.
-Udi Manber, Chief Scientist, Yahoo! (2001)

Coming up

- Bit vectors
- Non-dictionary data structures
- **Project 1** will be due next Tuesday
 - Sorting algorithms, Big-Oh analysis
- **Exam 1** will be returned Thursday
- **Recommended readings:** Sections 3.8-3.9
- **Practice problems:** 1-2 problems from "Trees and Other Dictionary Structures"