

Advanced data structures

William Hendrix

Outline

- Project overview
- Review
- Other dictionary implementations
 - Bit vectors
- Heaps
- Union-find
- Prefix and suffix trees (not going to cover)

Project 1

- Implement 4 sorting algorithms
 - SelectionSort
 - InsertionSort
 - MergeSort
 - QuickSort
 - Submit your code (C++ or Java)
- Run these algorithms with 4 kinds of input
 - Increasing, decreasing, random, and constant arrays
 - Array size 10k-100k
 - Submit table of results (10 x 16 + headers)
- Apply OLS regression to estimate time complexity
- Compare empirical complexity to theoretical (given in assignment)
 - Submit PDF with equations and analysis

Hash tables

- Apply hash function to map value into an allocated array
- Use one of two strategies to handle collisions
- Separate chaining
 - Each location is the head of a linked list
 - Append new element to list
 - Never "need" to reallocate
- Open addressing
 - Find the next open location, insert there
 - Can scan quadratically to avoid "congestion"
 - No links, so table can be larger with same memory
 - Linear scanning benefits from caching
 - Deleting an element requires reinserting everything that follows
- Expected-case complexity
 - Collisions are relatively infrequent ($O(1)$)
 - Array size (m) is $O(n)$

Hash table complexity

Operation	Hash tables worst-case	Hash tables expected-case	Balanced BST
Search(x)	$O(n)$	$O(1)$	$O(\lg n)$
Delete(x)	$O(n)/O(n^2)$	$O(1)$	$O(\lg n)$
Insert(x)	$O(n)$	$O(1)$	$O(\lg n)$
Build	$O(m + n^2)$	$O(n)$	$O(n \lg n)$
Resize	$O(m + n^2)$	$O(1)$, <i>amortized</i>	n/a
Min()	$O(m + n)/O(m)$	$O(n)$	$O(\lg n)$
Max()	$O(m + n)/O(m)$	$O(n)$	$O(\lg n)$
Predecessor(x)	$O(m + n)/O(m)$	$O(n)$	$O(\lg n)$
Successor(x)	$O(m + n)/O(m)$	$O(n)$	$O(\lg n)$

- This is amazing!
- The three most important techniques are hashing, hashing, and hashing.
-Udi Manber, Chief Scientist, Yahoo! (2001)

Hash tables: summary

- Poor worst-case complexity
- Excellent expected-case complexity
- Often fastest data structure in practice
 - Not as space-efficient as array
- **However:** be careful about expected-case assumptions
 - Hash function matters *intensely* for good performance
 - If too many values are mapped to the same location, we get worst-case performance
 - If data distribution includes lots of values that hash function collides, we get worst-case performance
 - Must be fast (impacts all operations)
 - Data distribution is also important
 - Some datasets are very skewed in frequency
 - Another data structure might be more appropriate
 - Need to ensure that we don't insert too many elements into table
 - Load factor: n / m
 - Usually resize above a certain threshold (e.g., 0.5, 0.75)

Bit vectors

- Represents a set as a sequence of Boolean values (bits)
 - One bit per value the set could contain
- **Example:**
 - Set of 1-100

00000000	00000000	00000000	00000000	00000000	00000000	00000000
8	16	24	32	40	48	56
00000000	00000000	00000000	00000000	00000000	0000	
64	72	80	88	96	100	

- Insert 3, 15, 25, 27, 82, 96, 100

Bit vectors

- Represents a set as a sequence of Boolean values (bits)
 - One bit per value the set could contain
- **Example:**
 - Set of 1-100

00100000	00000010	00000000	10100000	00000000	00000000	00000000
8	16	24	32	40	48	56
00000000	00000000	00000000	01000000	00000001	0001	
64	72	80	88	96	100	

- Insert 3, 15, 25, 27, 82, 96, 100
 - Byte = $\text{floor}((x - \text{min}) / 8)$
 - Bit = $(x - \text{min}) \bmod 8$
- Implemented as `char` array (C, C++, Java)
- Use bitwise operations
 - Left shift and right shift (\ll , \gg)
 - Bitwise and, or, xor, not ($\&$, $|$, \wedge , \sim)

Bit vector operations

- **Search(x)**
 - Test bit with bitwise and
 - `arr[byte] & (1 << bit)`
 - $O(1)$ time
- **Insert(x)**
 - Set bit with bitwise or
 - `arr[byte] |= (1 << bit)`
 - $O(1)$ time
- **Delete(x)**
 - Unset bit with bitwise and
 - `arr[byte] &= ~(1 << bit)`
 - $O(1)$ time

Summary: bit vector

Operation	Bit vector	Hash table (expected)	Balanced BST
Search(x)	$O(1)$	$O(1)$	$O(\lg n)$
Delete(x)	$O(1)$	$O(1)$	$O(\lg n)$
Insert(x)	$O(1)$	$O(1)$	$O(\lg n)$


- Constant time operations, even in worst case
- Very low coefficients
- Set operations (union, intersection, difference) can be implemented with bitwise ops
- Can only store integer data
- Cannot store duplicate values
- Space (and initialization/set op cost) determined by data *range*, not by number of values (n)
 - Inefficient if range not bounded

Priority queues

- Data structure for finding maxima or minima
 - E.g., in a greedy algorithm
- Main operations
 - **Insert(x)**: adds a value to the heap
 - **Max()/Min()**: returns max/min value
 - **DeleteMax()/DeleteMin()**: deletes the max/min value from the heap
- Primary implementation: heap
 - Complete binary tree (possibly incomplete bottom level)
 - Heap property
 - Every child has value \leq its parent's value (max-heap)
 - Every child has value \geq its parent's value (min-heap)
 - Consequence: root node is max (min)
- Extra op:
 - **Heapify(arr)**: builds a heap from an unsorted array

Heap implementation

- Implemented as array
- Root stored at index 1
- Children of element i stored at $2i$ and $2i+1$
 - Parent at $\text{floor}(i/2)$



- Because heaps are complete, array has no "gaps"

Priority queue operations

- Descriptions assume max-heap
- **Insert(x):**
 - Append to array
 - Swap with parents until parent is larger (or at root)
 - $O(\lg n)$ time
- **Max():**
 - Return root
 - $O(1)$ time

Priority queue deletion

- **DeleteMax()**
 - Swap max with last element and call PercolateDown(1)
 - **PercolateDown(i)**
 - If $heap[i]$ is smaller than its max child, swap
 - If so, repeat on that element until it is larger than max child or a leaf
 - $O(\lg n)$ time
 - $O(\lg n)$ time
- **Heapify()**
 - Call PercolateDown(i) from end to beginning
 - Half have no children, half of rest have 1 child, etc
 - $O(n)$ time total

```
1 Algorithm: Delete(x)
2 Swap  $heap[1]$  and  $heap[n]$ ;
3  $n = n - 1$ ;
4 PercolateDown(1);
```

```
1 Algorithm: PercolateDown(i)
2 if  $2i \leq n$  then
3    $mc = 2i$ ;
4   if  $mc + 1 \leq n$  and
      $heap[mc + 1] > heap[mc]$  then
5      $mc = mc + 1$ ;
6   end
7   if  $heap[i] < heap[mc]$  then
8     Swap  $heap[i]$  and  $heap[mc]$ ;
9     PercolateDown( $mc$ );
10  end
11 end
```

```
1 Algorithm: Heapify(i)
2 for  $i = \lfloor n/2 \rfloor$  to 1 step  $-1$ 
3   do
4     PercolateDown( $i$ );
5 end
```

Priority queue implementations

Operation	Heap	Unsorted array	Sorted array	Balanced BST	Fibonacci heap
Insert(x)	$O(\lg n)$	$O(1)$	$O(n)$	$O(\lg n)$	$O(1)$
Max()	$O(1)$	$O(1)$	$O(1)$	$O(1)$	$O(1)$
DeleteMax()	$O(\lg n)$	$O(n)$	$O(1)$	$O(\lg n)$	$O(\lg n)$, amort.
Build	$O(n)$	$O(n)$	$O(n \lg n)$	$O(n \lg n)$	$O(n)$

- BST has similar complexity, but higher coefficients
- Great at finding max (or min)
- Other operations (min/max, search, predecessor, etc.) are not good
 - Min-max heap can do either, but is more complex
- Fibonacci heap has even better complexity
 - More complex, higher coefficients, less space efficient
 - Fairly slow unless data is quite large

Union-Find data structure

- A.k.a., disjoint set data structure
- Used to represent a *partition*
 - Larger set split into smaller sets with no overlap
 - E.g., clusters, connected components
- Primary operations
 - **Find(x)**
 - Return the partition ID for element x
 - All elements in the same partition must return the same value
 - IDs might not be consecutive
 - **Union(a, b)**
 - Join partitions containing a and b

Union-Find implementation

- Array contains element IDs
- Partition IDs are elements pointing to themselves
- Initially, all elements are isolated:

0	1	2	3	4	5	6	7
{0}	{1}	{2}	{3}	{4}	{5}	{6}	{7}

- First try:
 - Find(x)
 - Follow links until you hit a partition ID
 - Return partition ID
 - $O(n)$ time
 - Union(a, b)
 - Point Find(a) to b
 - $O(n)$ time

```
1 Algorithm: Find(x)
2 if unionfind[x] = x then
3   | return x;
4 else
5   | id =
   | Find(unionfind[x])
   | return id;
6 end
```

```
1 Algorithm: Union(a, b)
2 id = Find(a);
3 unionfind[id] = b;
```

Coming up

- Finish union-find
- Sorting algorithms
- **Project 1** is posted on Canvas
 - Sorting algorithms, Big-Oh analysis
- **Recommended readings:** Sections 4.2 and 4.5
- **Practice problems:** p. 100: 1-2 problems from “Applications of Tree Structures”, attempt a problem from “Interview Problems”